

admin1@414-14:~\$ start-all.sh

This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh

Starting namenodes on [localhost]

admin1@localhost's password:

localhost: Permission denied, please try again.

admin1@localhost's password:

localhost: starting namenode, logging to /home/admin1/hadoop-2.7.3/logs/hadoop-admin1-namenode-414-14.out

admin1@localhost's password:

localhost: starting datanode, logging to /home/admin1/hadoop-2.7.3/logs/hadoop-admin1-datanode-414-14.out

Starting secondary namenodes [0.0.0.0]

admin1@0.0.0.0's password:

0.0.0.0: starting secondarynamenode, logging to /home/admin1/hadoop-2.7.3/logs/hadoop-admin1-secondarynamenode-414-14.out

starting yarn daemons

starting resourcemanager, logging to /home/admin1/hadoop-2.7.3/logs/yarn-admin1-resourcemanager-414-14.out

admin1@localhost's password:

localhost: starting nodemanager, logging to /home/admin1/hadoop-2.7.3/logs/yarn-admin1-nodemanager-414-14.out

admin1@414-14:~\$ jps

3107 NameNode

3735 ResourceManager

3306 DataNode

4283 Jps

4091 NodeManager

3550 SecondaryNameNode

admin1@414-14:~\$ mkdir TEAP

admin1@414-14:~\$ cd TEAP

admin1@414-14:~/TEAP\$ javac -classpath hadoop-core-1.2.1.jar Cyber.java

admin1@414-14:~/TEAP\$ cd ..

admin1@414-14:~\$ jar -cvf TEAP.jar -C TEAP/ .

added manifest

adding: Cyber.java(in = 2416) (out= 938)(deflated 61%)

adding: Cyber\$E_EMapper.class(in = 1949) (out= 813)(deflated 58%)

adding: Cyber.class(in = 1509) (out= 752)(deflated 50%)

adding: hadoop-core-1.2.1.jar(in = 4203713) (out= 3898883)(deflated 7%)

adding: Cyber\$E_EReducer.class(in = 1596) (out= 654)(deflated 59%)

admin1@414-14:~\$ hadoop fs -mkdir /TEP

admin1@414-14:~\$ hadoop fs -put sample.txt /TEP

admin1@414-14:~\$ hadoop jar TEAP.jar Cyber /TEP /TEO

22/04/18 11:51:23 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id

22/04/18 11:51:23 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker, sessionId=

22/04/18 11:51:23 INFO jvm.JvmMetrics: Cannot initialize JVM Metrics with processName=JobTracker, sessionId= - already initialized

22/04/18 11:51:23 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.

22/04/18 11:51:23 INFO mapred.FileInputFormat: Total input paths to process : 1

22/04/18 11:51:23 INFO mapreduce.JobSubmitter: number of splits:1

22/04/18 11:51:23 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local510926160_0001

22/04/18 11:51:24 INFO mapreduce.Job: The url to track the job: http://localhost:8080/

22/04/18 11:51:24 INFO mapred.LocalJobRunner: OutputCommitter set in config null

22/04/18 11:51:24 INFO mapreduce.Job: Running job: job_local510926160_0001

22/04/18 11:51:24 INFO mapred.LocalJobRunner: OutputCommitter is

org.apache.hadoop.mapred.FileOutputCommitter

22/04/18 11:51:24 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 1

22/04/18 11:51:24 INFO mapred.LocalJobRunner: Waiting for map tasks

22/04/18 11:51:24 INFO mapred.LocalJobRunner: Starting task: attempt_local510926160_0001_m_000000_0

22/04/18 11:51:24 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 1

22/04/18 11:51:24 INFO mapred.Task: Using ResourceCalculatorProcessTree : []

22/04/18 11:51:24 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/TEP/sample.txt:0+412

22/04/18 11:51:24 INFO mapred.MapTask: numReduceTasks: 1
22/04/18 11:51:24 INFO mapred.MapTask: (EQUATOR) 0 kvi 26214396(104857584)
22/04/18 11:51:24 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
22/04/18 11:51:24 INFO mapred.MapTask: soft limit at 83886080
22/04/18 11:51:24 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
22/04/18 11:51:24 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
22/04/18 11:51:24 INFO mapred.MapTask: Map output collector class =
org.apache.hadoop.mapred.MapTask\$MapOutputBuffer
22/04/18 11:51:24 INFO mapred.LocalJobRunner:
22/04/18 11:51:24 INFO mapred.MapTask: Starting flush of map output
22/04/18 11:51:24 INFO mapred.MapTask: Spilling map output
22/04/18 11:51:24 INFO mapred.MapTask: bufstart = 0; bufend = 206; bufvoid = 104857600
22/04/18 11:51:24 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26214320(104857280); length =
77/6553600
22/04/18 11:51:24 INFO mapred.MapTask: Finished spill 0
22/04/18 11:51:24 INFO mapred.Task: Task:attempt_local510926160_0001_m_000000_0 is done. And is in the
process of committing
22/04/18 11:51:24 INFO mapred.LocalJobRunner: hdfs://localhost:9000/TEP/sample.txt:0+412
22/04/18 11:51:24 INFO mapred.Task: Task 'attempt_local510926160_0001_m_000000_0' done.
22/04/18 11:51:24 INFO mapred.LocalJobRunner: Finishing task: attempt_local510926160_0001_m_000000_0
22/04/18 11:51:24 INFO mapred.LocalJobRunner: map task executor complete.
22/04/18 11:51:24 INFO mapred.LocalJobRunner: Waiting for reduce tasks
22/04/18 11:51:24 INFO mapred.LocalJobRunner: Starting task: attempt_local510926160_0001_r_000000_0
22/04/18 11:51:24 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 1
22/04/18 11:51:24 INFO mapred.Task: Using ResourceCalculatorProcessTree : []
22/04/18 11:51:24 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin:
org.apache.hadoop.mapreduce.task.reduce.Shuffle@754f32c6
22/04/18 11:51:24 INFO reduce.MergeManagerImpl: MergerManager: memoryLimit=334338464,
maxSingleShuffleLimit=83584616, mergeThreshold=220663392, ioSortFactor=10,
memToMemMergeOutputsThreshold=10
22/04/18 11:51:24 INFO reduce.EventFetcher: attempt_local510926160_0001_r_000000_0 Thread started:
EventFetcher for fetching Map Completion Events
22/04/18 11:51:24 INFO reduce.LocalFetcher: localfetcher#1 about to shuffle output of map
attempt_local510926160_0001_m_000000_0 decomp: 75 len: 79 to MEMORY
22/04/18 11:51:24 INFO reduce.InMemoryMapOutput: Read 75 bytes from map-output for
attempt_local510926160_0001_m_000000_0
22/04/18 11:51:24 INFO reduce.MergeManagerImpl: closeInMemoryFile -> map-output of size: 75,
inMemoryMapOutputs.size() -> 1, commitMemory -> 0, usedMemory -> 75
22/04/18 11:51:24 INFO reduce.EventFetcher: EventFetcher is interrupted.. Returning
22/04/18 11:51:24 INFO mapred.LocalJobRunner: 1 / 1 copied.
22/04/18 11:51:24 INFO reduce.MergeManagerImpl: finalMerge called with 1 in-memory map-outputs and 0 on-disk
map-outputs
22/04/18 11:51:24 INFO mapred.Merger: Merging 1 sorted segments
22/04/18 11:51:24 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total size: 67 bytes
22/04/18 11:51:24 INFO reduce.MergeManagerImpl: Merged 1 segments, 75 bytes to disk to satisfy reduce memory
limit
22/04/18 11:51:24 INFO reduce.MergeManagerImpl: Merging 1 files, 79 bytes from disk
22/04/18 11:51:24 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce
22/04/18 11:51:24 INFO mapred.Merger: Merging 1 sorted segments
22/04/18 11:51:24 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total size: 67 bytes
22/04/18 11:51:24 INFO mapred.LocalJobRunner: 1 / 1 copied.
22/04/18 11:51:24 INFO mapred.Task: Task:attempt_local510926160_0001_r_000000_0 is done. And is in the process
of committing
22/04/18 11:51:24 INFO mapred.LocalJobRunner: 1 / 1 copied.
22/04/18 11:51:24 INFO mapred.Task: Task attempt_local510926160_0001_r_000000_0 is allowed to commit now
22/04/18 11:51:24 INFO output.FileOutputCommitter: Saved output of task
'attempt_local510926160_0001_r_000000_0' to
hdfs://localhost:9000/TEO/_temporary/0/task_local510926160_0001_r_000000
22/04/18 11:51:24 INFO mapred.LocalJobRunner: reduce > reduce

22/04/18 11:51:24 INFO mapred.Task: Task 'attempt_local510926160_0001_r_000000_0' done.
22/04/18 11:51:24 INFO mapred.LocalJobRunner: Finishing task: attempt_local510926160_0001_r_000000_0
22/04/18 11:51:24 INFO mapred.LocalJobRunner: reduce task executor complete.
22/04/18 11:51:25 INFO mapreduce.Job: Job job_local510926160_0001 running in uber mode : false
22/04/18 11:51:25 INFO mapreduce.Job: map 100% reduce 100%
22/04/18 11:51:25 INFO mapreduce.Job: Job job_local510926160_0001 completed successfully
22/04/18 11:51:25 INFO mapreduce.Job: Counters: 35

File System Counters

FILE: Number of bytes read=7806492
FILE: Number of bytes written=8477739
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=824
HDFS: Number of bytes written=49
HDFS: Number of read operations=13
HDFS: Number of large read operations=0
HDFS: Number of write operations=4

Map-Reduce Framework

Map input records=20
Map output records=20
Map output bytes=206
Map output materialized bytes=79
Input split bytes=88
Combine input records=20
Combine output records=6
Reduce input groups=6
Reduce shuffle bytes=79
Reduce input records=6
Reduce output records=6
Spilled Records=12
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=0
Total committed heap usage (bytes)=530579456

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=412

File Output Format Counters

Bytes Written=49

admin1@414-14:~\$ hadoop fs -ls /TEO

Found 2 items

```
-rw-r--r--  1 admin1 supergroup      0 2022-04-18 11:51 /TEO/_SUCCESS
-rw-r--r--  1 admin1 supergroup    49 2022-04-18 11:51 /TEO/part-00000
```

admin1@414-14:~\$ hadoop fs -cat /TEO/part-00000

Ashok 6
Heena 6
Kunda 8
Pratap 7
Raman 7
Rohan 6

admin1@414-14:~\$ ^C
admin1@414-14:~\$

OUTPUT:

```
File Edit View Search Terminal Help
    Reduce output records=6
    Spilled Records=12
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=0
    Total committed heap usage (bytes)=530579456
Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
File Input Format Counters
    Bytes Read=412
File Output Format Counters
    Bytes Written=49
admin1@414-14:~$ hadoop fs -ls /TEO
Found 2 items
-rw-r--r--  1 admin1 supergroup          0 2022-04-18 11:51 /TEO/_SUCCESS
-rw-r--r--  1 admin1 supergroup        49 2022-04-18 11:51 /TEO/part-00000
admin1@414-14:~$ hadoop fs -cat /TEO/part-r-00000
cat: '/TEO/part-r-00000': No such file or directory
admin1@414-14:~$ hadoop fs -cat /TEO/part-00000
Ashok      6
Heena      6
Kunda      8
Pratap     7
Raman      7
Rohan      6
admin1@414-14:~$ ^C
admin1@414-14:~$
```